

Harish Manoharan

F-1 OPT (May 2026) | +1 (945) 216-8488 | harish_manoharan@outlook.com | linkedin.com/in/harishm17 | [harishm17.github.io](https://github.com/harishm17)

EXPERIENCE

Software Engineer (Core Team) <i>Purgo AI</i>	Jun 2025 - Present San Francisco Bay Area
<ul style="list-style-type: none">Engineered a full-stack validation platform for Databricks, automating IQ/OQ validation workflows to reduce validation cycles from days to 20 minutes via orchestrated API testing and audit-ready PDF generationDesigned and implemented an LLM Evaluation Service, automating 20+ reliability checks (hallucination rate, PII leakage, traceability) to verify model safety and support enterprise audit readinessBuilt a graph-based retrieval layer that indexes Databricks catalog schemas and table lineage, serving semantic context to the code generation agent to boost SQL/PySpark answer correctness by 12%Implemented an agent orchestration state machine (LangGraph Supervisor-Worker) that routes tasks between specialized agents (Research, Design, Coding) to reliably execute end-to-end data engineering workflowsArchitected an embedded eSignature workflow (DocuSign API, OAuth2, webhooks), enabling in-platform signing and cutting approval cycle time by 50% through server-side envelope creation and event-driven status updates	
Software Engineer Intern <i>Digitus Business Solutions</i>	May 2024 - Jul 2024 Bangalore, India
<ul style="list-style-type: none">Engineered a text-to-SQL agent with dynamic schema discovery and execution-feedback retries, improving SQL execution accuracy from 52% to 76% on a 500-query evaluation set by automating joins across complex schemasBuilt a multimodal RAG system indexing internal knowledge (PPTs, videos) with layout-aware chunking and citation enforcement, reducing median time-to-resolution (TTR) from 20 min to 6 minDelivered a human-in-the-loop contract tool with risk detection, shortening reviews from 5 days to 2 days	
Technology Consulting Intern - Data Analytics <i>PricewaterhouseCoopers</i>	May 2023 - Jul 2023 Bangalore, India
<ul style="list-style-type: none">Built a sales forecasting pipeline using an XGBoost-SARIMA ensemble with feature engineering for seasonality and promotions, reducing MAPE by 15% over the baseline model for 50+ product SKUsEngineered an aspect-based sentiment analysis pipeline on 30k+ reviews to identify key satisfaction drivers	

PROJECTS

Realtime Exam Prep Coach	Live GitHub
<ul style="list-style-type: none">Built an AI exam-prep platform that converts course materials into topic-structured notes, quizzes, and graded practice exams using a multi-stage RAG pipeline with semantic chunking and hybrid retrievalImplemented a real-time voice coach using WebRTC + OpenAI Realtime with tool-driven drills, concept-only guardrails, automated scoring, and latency/quality telemetry; achieved p95 time-to-first-audio of 600 msDeployed on GCP (Cloud Run/Cloud Tasks) with GitHub Actions CI/CD, Prisma + Postgres, and Google OAuth	
Multi-agent Email Assistant	GitHub
<ul style="list-style-type: none">Built a tool-orchestrated email and calendar copilot that answers questions over inbox and schedule, drafts reply with personalized writing style, and manages calendar events through Google Workspace integrationBuilt a safety and reliability stack with schema-validated tool outputs, PII redaction, policy checks, and an independent LLM judge; added robust fallbacks and observability for latency and tool failures	
Unsupervised Behavior Discovery and Genetics of Mice - Master's Thesis (with The Jackson Laboratory)	
<ul style="list-style-type: none">Temporal Video segmentation and action recognition using key point data from large mice video datasetTrained a bidirectional RNN encoder-decoder model and performed k-means clustering to segment videos	

TECHNICAL SKILLS

Programming Languages: Python, C++, TypeScript/JavaScript, SQL

LLM Systems: RAG, AI Agents, LangGraph/LangChain, PyTorch, Langfuse (tracing), LLM evaluation, guardrails/safety

Cloud, Data and Backend: API design, AWS/GCP, Docker, FastAPI, Node.js, Postgres, PySpark, MongoDB, Redis

EDUCATION

M.S. in Computer Science <i>University of Texas at Dallas - Dean's Graduate Scholar</i>	Expected May 2026 GPA: 3.91/4.0
Interdisciplinary Dual Degree: M.Tech, Data Science & B.S., Biological Sciences <i>Indian Institute of Technology Madras - INSPIRE Scholar, Young Research Fellow</i>	May 2024 GPA: 8.4/10.0